


```
1 import pandas as pd
2 import numpy as np
3 df=pd.read_excel('/content/drive/MyDrive/Datasets/Road.xlsx')
4 df
```



	Sex_of_driver	Educational_level	Vehicle_driver_relation	Owner_of_vehicle	Area
0	Male	Above high school	Employee	Owner	
1	Male	Junior high school	Employee	Owner	
2	Male	Junior high school	Employee	Owner	
3	Male	Junior high school	Employee	Governmental	
4	Male	Junior high school	Employee	Owner	
...	...	...	...	...	...
12311	Male	NaN	Employee	Owner	
12312	Male	Elementary school	Employee	Owner	
12313	Male	Junior high school	Employee	Owner	
12314	Female	Junior high school	Employee	Owner	
12315	Male	Junior high school	Employee	Owner	


12316 rows × 12 columns

Next steps:

Generate code with df

 View recommended plots

```
1 df.head()
```



	Sex_of_driver	Educational_level	Vehicle_driver_relation	Owner_of_vehicle	Area_acc
0	Male	Above high school	Employee	Owner	f
1	Male	Junior high school	Employee	Owner	
2	Male	Junior high school	Employee	Owner	Re
3	Male	Junior high school	Employee	Governmental	
4	Male	Junior high school	Employee	Owner	

Next steps:


Generate code with df

 View recommended plots

```
1 df.replace('Unknown',np.NaN,inplace=True)


1 df.replace('Fatal injury','Serious Injury',inplace=True)

1 df['Accident_severity'].value_counts()
```



Accident_severity	
Slight Injury	10415
Serious Injury	1901
Name: count, dtype: int64	

```
1 df.isna().sum()
```




Sex_of_driver	178
Educational_level	841
Vehicle_driver_relation	593
Owner_of_vehicle	482
Area_accident_occured	261
Light_conditions	0
Weather_conditions	292

```
Number_of_vehicles_involved    0
Number_of_casualties           0
Vehicle_movement               396
Cause_of_accident              25
Accident_severity              0
dtype: int64
```


```
1 df.dropna(inplace=True)
```

```
1 df.describe()
```




	Number_of_vehicles_involved	Number_of_casualties
count	10064.000000	10064.000000
mean	2.042130	1.546999
std	0.694532	1.010320
min	1.000000	1.000000
25%	2.000000	1.000000
50%	2.000000	1.000000
75%	2.000000	2.000000
max	7.000000	8.000000

```
1 df.info()
```




```
<class 'pandas.core.frame.DataFrame'>
Index: 10064 entries, 0 to 12315
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Sex_of_driver                         10064 non-null  object
1   Educational_level                     10064 non-null  object
2   Vehicle_driver_relation                10064 non-null  object
3   Owner_of_vehicle                      10064 non-null  object
4   Area_accident_occured                 10064 non-null  object
5   Light_conditions                      10064 non-null  object
6   Weather_conditions                    10064 non-null  object
7   Number_of_vehicles_involved           10064 non-null  int64
8   Number_of_casualties                  10064 non-null  int64
9   Vehicle_movement                      10064 non-null  object
10  Cause_of_accident                     10064 non-null  object
11  Accident_severity                     10064 non-null  object
dtypes: int64(2), object(10)
memory usage: 1022.1+ KB
```

```
1 df.shape
```



```
(10064, 12)
```


```
1 df.dtypes
```



```
Sex_of_driver                object
Educational_level            object
Vehicle_driver_relation       object
Owner_of_vehicle              object
Area_accident_occured         object
Light_conditions              object
Weather_conditions            object
Number_of_vehicles_involved   int64
Number_of_casualties          int64
Vehicle_movement              object
Cause_of_accident             object
Accident_severity             object
dtype: object
```

```
1 df.reset_index(drop=True,inplace=True)
```

```
1 df
```



	Sex_of_driver	Educational_level	Vehicle_driver_relation	Owner_of_vehicle	Area
0	Male	Above high school	Employee	Owner	
1	Male	Junior high school	Employee	Owner	
2	Male	Junior high school	Employee	Owner	
3	Male	Junior high school	Employee	Governmental	
4	Male	Junior high school	Employee	Owner	
...	...	...	...	...	...
10059	Female	Elementary school	Employee	Owner	
10060	Male	Elementary school	Employee	Owner	
10061	Male	Junior high school	Employee	Owner	
10062	Female	Junior high school	Employee	Owner	
10063	Male	Junior high school	Employee	Owner	


10064 rows × 12 columns

Next steps:

Generate code with df


 View recommended plots

```
1 df.dtypes
```



Sex_of_driver	object
Educational_level	object
Vehicle_driver_relation	object
Owner_of_vehicle	object
Area_accident_occured	object
Light_conditions	object
Weather_conditions	object
Number_of_vehicles_involved	int64
Number_of_casualties	int64
Vehicle_movement	object
Cause_of_accident	object
Accident_severity	object
dtype:	object


```
1 df.shape
```



(10064, 12)

## ✓ Label Encoder

```
1 for i in df:
2     if df[i].dtypes=='object':
3         print(i,'|',df[i].unique())
```



Sex_of_driver		['Male' 'Female']
Educational_level		['Above high school' 'Junior high school' 'Elementary school' 'High school' 'Illiterate' 'Writing & reading']
Vehicle_driver_relation		['Employee' 'Owner' 'Other']
Owner_of_vehicle		['Owner' 'Governmental' 'Organization' 'Other']
Area_accident_occured		['Residential areas' 'Office areas' 'Recreational areas' 'Industrial areas' 'Other' 'Church areas' 'Market areas' 'Rural village areas' 'Outside rural areas' 'Hospital areas' 'School areas' 'Rural village areasOffice areas']
Light_conditions		['Daylight' 'Darkness - lights lit' 'Darkness - no lighting' 'Darkness - lights unlit']
Weather_conditions		['Normal' 'Raining' 'Raining and Windy' 'Cloudy' 'Other' 'Windy' 'Snow' 'Fog or mist']
Vehicle_movement		['Going straight' 'Moving Backward' 'U-Turn' 'Turnover' 'Waiting to go' 'Getting off' 'Reversing' 'Parked' 'Stopping' 'Other' 'Overtaking' 'Entering a junction']
Cause_of_accident		['Moving Backward' 'Overtaking' 'Changing lane to the left' 'Changing lane to the right' 'Other' 'No priority to vehicle' 'No priority to pedestrian' 'No distancing']

```

'Getting off the vehicle improperly' 'Overloading' 'Driving carelessly'
'Driving at high speed' 'Driving to the left' 'Overturning' 'Turnover'
'Driving under the influence of drugs' 'Overspeed' 'Drunk driving'
'Improper parking']
Accident_severity | ['Slight Injury' 'Serious Injury']

```

```

1 from sklearn.preprocessing import LabelEncoder
2 lst=['Sex_of_driver','Educational_level','Vehicle_driver_relation','Owner_of_vehicle','Area_accident_occured','Light_conditions','Weather
3 label={}
4 for col in lst:
5     label[col]=LabelEncoder()
6     df[col]=label[col].fit_transform(df[col])
7 label

```

```

↻ {'Sex_of_driver': LabelEncoder(),
   'Educational_level': LabelEncoder(),
   'Vehicle_driver_relation': LabelEncoder(),
   'Owner_of_vehicle': LabelEncoder(),
   'Area_accident_occured': LabelEncoder(),
   'Light_conditions': LabelEncoder(),
   'Weather_conditions': LabelEncoder(),
   'Vehicle_movement': LabelEncoder(),
   'Cause_of_accident': LabelEncoder(),
   'Accident_severity': LabelEncoder()}

```

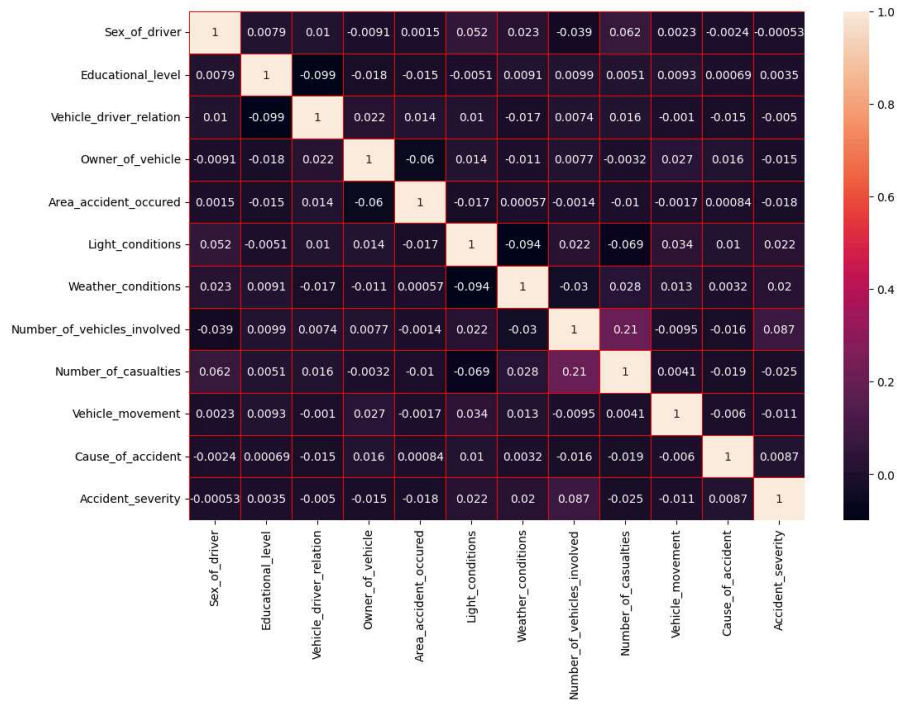
## Visualization

### ✓ Heatmap

```


1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 plt.subplots(figsize=(12, 8))
4 sns.heatmap(df.corr(),linewidths=0.5,linecolor='red',annot=True)
5 plt.show()

```



## Pair Plot


```
1 sns.pairplot(df,hue='Accident_severity',kind='scatter')
2 plt.show()
```



✓ Seaborn Plot

```
1 col=df[['Sex_of_driver','Educational_level','Vehicle_driver_relation','Owner_of_vehicle','Area_accident_occured','Light_conditions','Weat
2 sns.boxplot(col)
```

↗ <Axes: >



## Remove Outliers

```
1 lst=['Sex_of_driver','Educational_level','Vehicle_driver_relation','Owner_of_vehicle','Area_accident_occured','Light_conditions','Weather
2 for i in lst:
3     if df[i].dtype in ['int64']:
4         Q1 = df[i].quantile(0.25)
5         Q3 = df[i].quantile(0.75)
6         IQR = Q3 - Q1
7         lower_bound = Q1 - 1.5 * IQR
8         upper_bound = Q3 + 1.5 * IQR
9         df = df[(df[i] >= lower_bound) & (df[i] <= upper_bound)]
```

```
1 df['Accident_severity'].value_counts()
```


↗ Accident\_severity

1	1815
0	264

Name: count, dtype: int64

## Hist Plot

```
1 plt.hist(df['Accident_severity'])
2 plt.show()
```



```
1 df['Accident_severity'].value_counts()
```

```

↳ Accident_severity
1    1815
0     264
Name: count, dtype: int64

```

```

1 x=df.drop(columns='Accident_severity')
2 y=df['Accident_severity']

```

## ✓ OverSampling

```

1 from imblearn.over_sampling import RandomOverSampler
2 rand=RandomOverSampler()
3 x,y=rand.fit_resample(x,y)

```

```
1 df['Accident_severity'].value_counts()
```

```

↳ Accident_severity
1    1815
0     264
Name: count, dtype: int64

```

## ✓ TrainTest Splitting

```

1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)

```

## ✓ Normalized Scaling

```

1 from sklearn.preprocessing import StandardScaler
2 ss=StandardScaler()
3 x_train_scaled=ss.fit_transform(x_train)
4 x_test_scaled=ss.transform(x_test)

```

## ✓ KNeighborsClassifier

```

1 from sklearn.neighbors import KNeighborsClassifier
2 knn=KNeighborsClassifier()
3 knn.fit(x_train_scaled,y_train)

```



↻ ▾ KNeighborsClassifier  
KNeighborsClassifier()

```
1 y_pred=knn.predict(x_test_scaled)
2 y_pred
```

↻ array([0, 1, 0, ..., 1, 1, 0])

```
1 y_test.values
```

↻ array([1, 0, 1, ..., 1, 1, 0])

```
1 from sklearn.metrics import classification_report,accuracy_score
2 print(accuracy_score(y_test,y_pred)*100)
```

↻ 67.76859504132231

```
1 print(classification_report(y_test,y_pred))
```

↻

	precision	recall	f1-score	support
0	0.63	0.80	0.70	522
1	0.76	0.56	0.65	567
accuracy			0.68	1089
macro avg	0.69	0.68	0.67	1089
weighted avg	0.69	0.68	0.67	1089

```
1 x_scaled=ss.transform(x)
```

## ✓ SVM

```
1 from sklearn.svm import SVC
2 model=SVC()
3 model.fit(x_train_scaled,y_train)
```

↻ ▾ SVC  
SVC()

```
1 y_pred=model.predict(x_test_scaled)
2 y_pred
```

↻ array([0, 1, 0, ..., 1, 1, 0])

```
1 print(accuracy_score(y_test,y_pred)*100)
```

↻ 59.87144168962351

```
1 print(classification_report(y_test,y_pred))
```

↻

	precision	recall	f1-score	support
0	0.57	0.69	0.62	522
1	0.64	0.51	0.57	567
accuracy			0.60	1089
macro avg	0.61	0.60	0.60	1089
weighted avg	0.61	0.60	0.60	1089

## ✓ Naive Bayes

```

1 from sklearn.naive_bayes import MultinomialNB,GaussianNB,BernoulliNB
2 model1=BernoulliNB()
3 model1.fit(x_train_scaled,y_train)

```

↗

▾ BernoulliNB  
 BernoulliNB()

```

1 y_pred=model1.predict(x_test_scaled)

1 print(accuracy_score(y_test,y_pred)*100)

```

↗ 53.168044077134994

```

1 print(classification_report(y_test,y_pred))

```

↗

	precision	recall	f1-score	support
0	0.51	0.83	0.63	522
1	0.62	0.25	0.36	567
accuracy			0.53	1089
macro avg	0.57	0.54	0.50	1089
weighted avg	0.57	0.53	0.49	1089

## ▾ Decision Tree

```

1 from sklearn.tree import DecisionTreeClassifier,plot_tree
2 model2=DecisionTreeClassifier()
3 model2.fit(x_train,y_train)

```

↗

▾ DecisionTreeClassifier  
 DecisionTreeClassifier()

```

1 y_pred=model2.predict(x_test)
2 y_pred

```

↗ array([1, 1, 1, ..., 1, 1, 0])

```

1 print(accuracy_score(y_test,y_pred)*100)

```

↗ 73.27823691460054

```

1 print(classification_report(y_test,y_pred))

```

↗


	precision	recall	f1-score	support
0	0.66	0.90	0.76	522
1	0.86	0.58	0.69	567
accuracy			0.73	1089
macro avg	0.76	0.74	0.73	1089
weighted avg	0.77	0.73	0.73	1089

## ▾ Decision Tree Plot

```

1 plt.figure(figsize=(10,8))
2 plot_tree(model2,filled=True,feature_names=x_train.columns)
3 plt.show()

```



## RandomForest

```
1 from sklearn.ensemble import RandomForestClassifier
2 rf=RandomForestClassifier()
3 rf.fit(x_train,y_train)
```



```
▼ RandomForestClassifier
RandomForestClassifier()
```

```
1 y_pred=rf.predict(x_test)
2 y_pred
```



```
array([1, 1, 1, ..., 1, 1, 0])
```

```
1 print(accuracy_score(y_test,y_pred)*100)
```



```
73.46189164370982
```

```
1 print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.67	0.89	0.76	522
1	0.85	0.59	0.70	567
accuracy			0.73	1089
macro avg	0.76	0.74	0.73	1089
weighted avg	0.76	0.73	0.73	1089

## GradientBoosting

```
1 from sklearn.ensemble import GradientBoostingClassifier
2 gb=GradientBoostingClassifier()
3 gb.fit(x_train,y_train)
```

```
gc = GradientBoostingClassifier()
```

```
↳ GradientBoostingClassifier
GradientBoostingClassifier()
```

```
1 y_pred=gc.predict(x_test)
2 y_pred
```

```
↳ array([1, 1, 0, ..., 1, 1, 0])
```

```
1 print(accuracy_score(y_test,y_pred)*100)
```

```
↳ 61.89164370982553
```

```
1 print(classification_report(y_test,y_pred))
```

```
↳
```

	precision	recall	f1-score	support
0	0.58	0.73	0.65	522
1	0.68	0.52	0.59	567
accuracy			0.62	1089
macro avg	0.63	0.62	0.62	1089
weighted avg	0.63	0.62	0.62	1089

## ✖ XGBooster

```
1 from xgboost import XGBClassifier
2 xg=XGBClassifier()
3 xg.fit(x_train,y_train)
```

```
↳ XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...)
```

```
1 y_pred=gc.predict(x_test)
2 y_pred
```

```
↳ array([1, 1, 0, ..., 1, 1, 0])
```

```
1 print(accuracy_score(y_test,y_pred)*100)
```

```
↳ 61.89164370982553
```

```
1 print(classification_report(y_test,y_pred))
```

```
↳
```

	precision	recall	f1-score	support
0	0.58	0.73	0.65	522
1	0.68	0.52	0.59	567
accuracy			0.62	1089
macro avg	0.63	0.62	0.62	1089
weighted avg	0.63	0.62	0.62	1089

## ✖ AdaBooster

```
1 from sklearn.ensemble import AdaBoostClassifier
```

```
2 ada=AdaBoostClassifier()
3 ada.fit(x_train,y_train)
```

↗

▾ AdaBoostClassifier  
 AdaBoostClassifier()

```
1 y_pred=ada.predict(x_test)
2 y_pred
```

↗ array([1, 1, 0, ..., 0, 0, 0])

```
1 print(accuracy_score(y_test,y_pred)*100)
```

↗ 56.84113865932048

```
1 print(classification_report(y_test,y_pred))
```

↗

	precision	recall	f1-score	support
0	0.54	0.63	0.58	522
1	0.60	0.51	0.55	567
accuracy			0.57	1089
macro avg	0.57	0.57	0.57	1089
weighted avg	0.57	0.57	0.57	1089

```
1 from sklearn.model_selection import RandomizedSearchCV
2 params={'criterion':['entropy','gini'],
3         'max_depth':[5,10,30,40,50],
4         'min_samples_split':[10,20,25,30,45,60,100] }
5 rs=RandomizedSearchCV(RandomForestClassifier(),params,cv=10,n_iter=15)
6 rs.fit(x_train,y_train)
```

↗

▸ RandomizedSearchCV  
 ▸ estimator: RandomForestClassifier  
 ▸ RandomForestClassifier

```
1 rs.best_params_
```

↗ {'min\_samples\_split': 20, 'max\_depth': 30, 'criterion': 'gini'}

```
1 rs.best_score_
```

↗ 0.7044681179558437

```
1 knn=KNeighborsClassifier()
2 svc=SVC()
3 naive=BernoulliNB()
4 decision=DecisionTreeClassifier()
5 rf=RandomForestClassifier()
6 gcv=GradientBoostingClassifier()
7 xgb=XGBClassifier()
8 ada=AdaBoostClassifier()
```

```
1 model_lst=[knn,svc,naive,decision,rf,gcv,xgb,ada]
2 for model in model_lst:
3     model.fit(x_train,y_train)
4     y_pred=model.predict(x_test)
5     print(model)
6     print('Accuracy_Score:',accuracy_score(y_test,y_pred)*100)
7     print(classification_report(y_test,y_pred))
8     print('-----')
```

↗

accuracy			0.73	1089
macro avg	0.76	0.74	0.73	1089
weighted avg	0.77	0.73	0.73	1089

-----

GradientBoostingClassifier()  
Accuracy\_Score: 61.89164370982553

	precision	recall	f1-score	support
0	0.58	0.73	0.65	522
1	0.68	0.52	0.59	567

accuracy			0.62	1089
macro avg	0.63	0.62	0.62	1089
weighted avg	0.63	0.62	0.62	1089

-----

XGBClassifier(base\_score=None, booster=None, callbacks=None, colsample\_bylevel=None, colsample\_bynode=None, colsample\_bytree=None, device=None, early\_stopping\_rounds=None, enable\_categorical=False, eval\_metric=None, feature\_types=None, gamma=None, grow\_policy=None, importance\_type=None, interaction\_constraints=None, learning\_rate=None, max\_bin=None, max\_cat\_threshold=None, max\_cat\_to\_onehot=None, max\_delta\_step=None, max\_depth=None, max\_leaves=None, min\_child\_weight=None, missing=nan, monotone\_constraints=None, multi\_strategy=None, n\_estimators=None, n\_jobs=None, num\_parallel\_tree=None, random\_state=None, ...)

Accuracy\_Score: 71.71717171717171

	precision	recall	f1-score	support
--	-----------	--------	----------	---------